


Booting & Shutting Down




Reference Material:

- [1] Unified Extensible Firmware Interface Specification (Sections 2 & 3) https://uefi.org/sites/default/files/resources/UEFI%20Spec%202_6.pdf
- [2] UEFI, Debian wiki: <https://wiki.debian.org/UEFI>
- [3] UNIX and Linux System Administration Handbook, Section 2 (Booting and System Management Daemons).
- [4] systemd, Archlinux wiki: <https://wiki.archlinux.org/index.php/systemd>


Index

- **Introduction**
- **Booting, Stage 1: Hardware**
- **Booting, Stage 2: Bootloader (GRUB)**
- **Booting, Stage 3: Kernel**
- **Booting, Stage 4: Systemd**
- **Shutting Down**




Introduction

- Booting/Shutting Down are complex procedures, but system provides mechanisms to deal with them.
- ... However, this is one of the potential troubles of administration.
- Goals of this Chapter:
 - Understand the basic operation of both procedures.
 - Being able to customize them.
 - Being able to solve generic problems related to Boot process.
- Before: BIOS (and MBR) + SysV (see appendix), **Now: UEFI (and GPT) + Systemd**
 - Be careful with your online searches...





Introduction

- The main target of Booting process is loading kernel (OS) in memory and starting its execution.
 - ¿Where is the kernel before booting?
 - ¿What's the content of memory before booting?
- It is a sequential process:



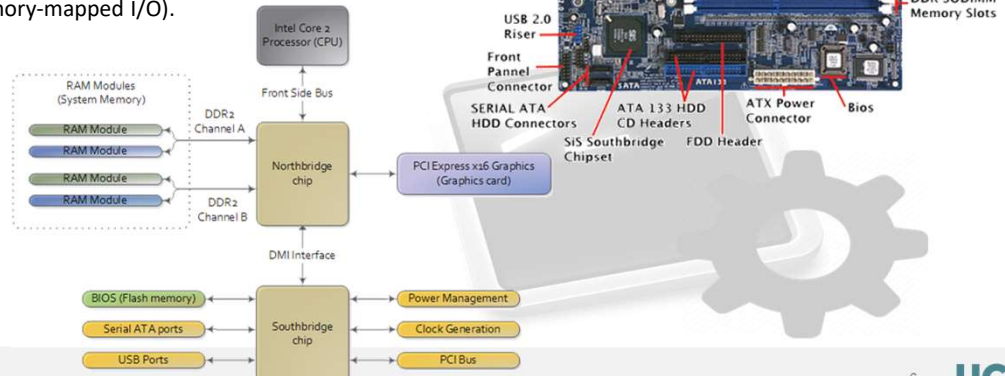
Index

- Introduction
- **Booting, Stage 1: Hardware+UEFI**
- Booting, Stage 2: Bootloader (GRUB)
- Booting, Stage 3: Kernel
- Booting, Stage 4: Systemd
- Shutting Down



Stage 1: Hardware

- Power-On:
 - After pushing Power-On button, the **Reset Vector** indicates the CPU the direction of the first instruction to execute (FFFFFFFF0h for x86). Such direction corresponds to an EPROM/Flash (motherboard) that stores the code corresponding to the Firmware (memory-mapped I/O).



Initialization Firmware


The old days...

- **BIOS (Basic Input Output System)**
 - firmware developed for the IBM PC/XT in the late 1970s
 - Adapted and expanded many times.
 - Support for MBR Partition Standard
 - 16-bit system, no disk driver available (no notion of files/directories in a disk).

Currently


- **UEFI: Unified Extensible Firmware Interface**
 - Firmware standard derived from Intel's **EFI** (2000).
 - Intel's alternative for its new Itanium processor
 - Better disk/network support: full support for **GUID Partition Table (GPT)** (see slides 10-11) and IPv6.
 - Driver to "understand" FAT filesystem (format of the EFI System Partition or ESP).
 - No real need for a bootloader (**no GRUB required**).
 - Improved Security. UEFI **Secureboot** (run only signed apps) to prevent pre-OS malware (bootkit).
 - Closer integration between OS and pre-boot environment: Requires support from the OS (Linux, OSX, Windows10).

Stage 1: Hardware + UEFI (Boot Manager)




- **Power-on-self-test (POST)**: examination, verification and start up of hardware devices (CPU, RAM, Controllers, etc.)
- Initialize the CPU with hardware-specific code (loaded directly from flash memory)
 - DRAM not yet available, UEFI firmware uses CPU cache as RAM.
- Initialize **Main memory**, as well as hardware components required for the next phases
- Load the **device drivers** from the flash memory into main memory (now available), initialize all required hardware (disk (FAT filesystem), monitor, keyboard,...) and register/use protocols.
 - Protocol: provide text output to the console or access to a PCI device.

Stage 1: Hardware + UEFI (Boot Manager)



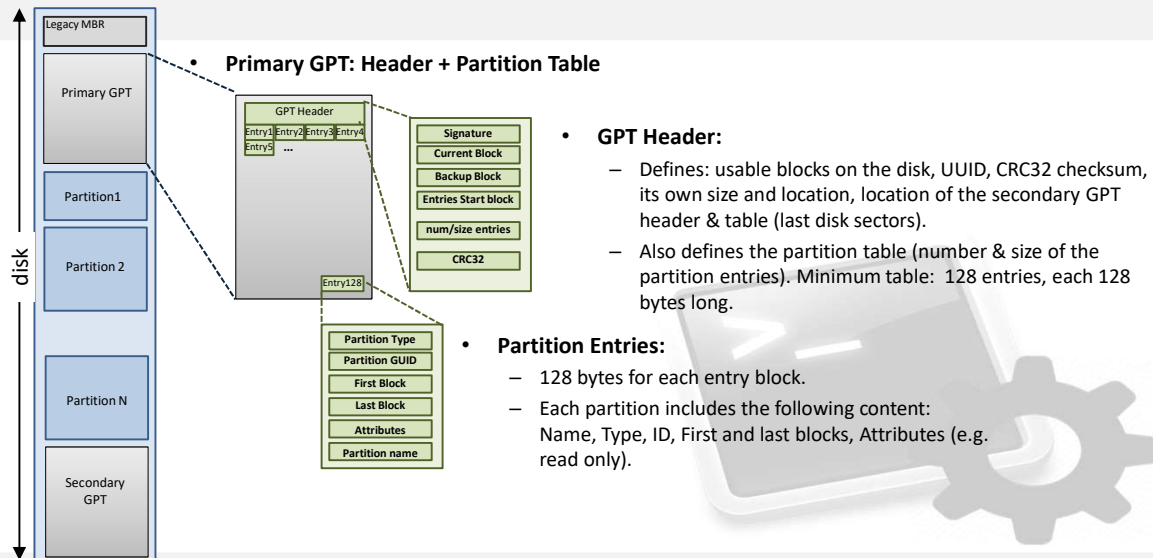
- **UEFI Boot Manager:** attempts to load UEFI applications (.efi files) in a predefined device order
 - Devices can be: storage, network.
 - Apps can be: **OS bootloader**, efi kernel, additional drivers (ext4), shell, GUI, etc.
 - **NVRAM** variables define this order
- Applications must reside on an UEFI-defined file system
 - **FAT** filesystem (format of the EFI System Partition or ESP).

Stage 1: Hardware + UEFI (Boot Manager)

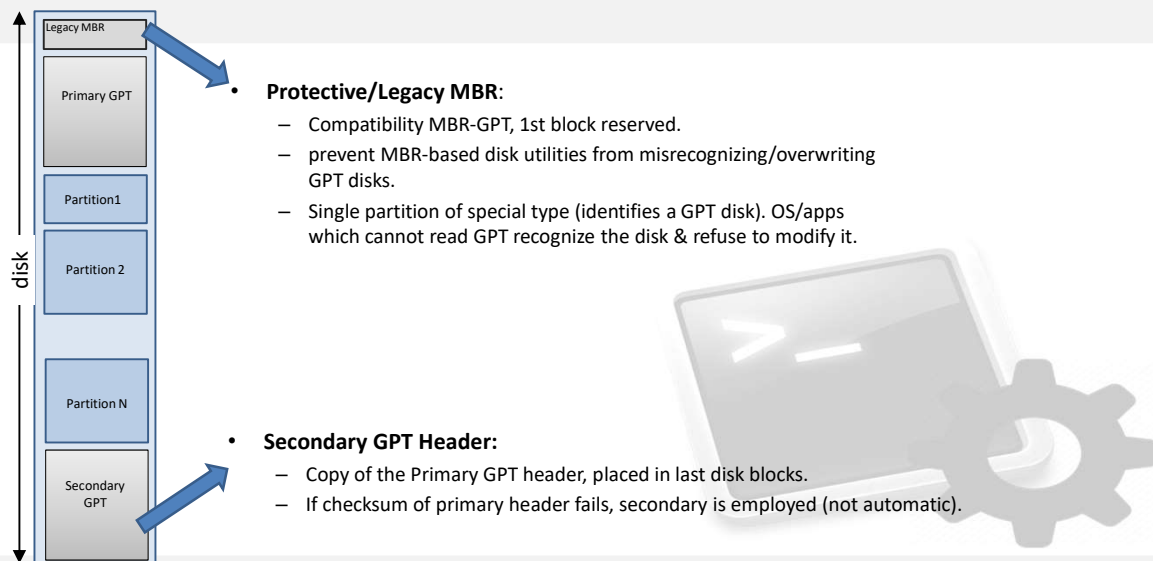


- **How is Apps & Loader search process performed?**
- UEFI consults the GPT partition table to identify the ESP.
 - Remember the partition creation process in gdisk (Partition Type)
 - It then reads the target application (.efi file) from a file in the ESP and executes it.
 - Pathname to load: configuration parameter. By default (Debian): /efi/boot/bootx64.efi
- Each installed OS has its own directory in EFI partition.
- If no bootloader is used (EFI stub support enabled in kernel), all files required to load the OS (kernel, ramdisk, etc.) must be available in this partition.


Remember: GPT Disks & Partitions



Remember: GPT Disks & Partitions



Stage 1: Hardware + UEFI (Boot Manager)



• The UEFI shell


- Keep this in mind, EFI is a small OS built into the motherboard, could have its own shell. Usually, internal part of EFI, if not, can be downloaded as any other app.
- Press “ESC” after Power On. Boot Manager -> EFI Internal Shell.
- Many commands similar to Linux shell commands:
 - cp, ls, rm, mv, touch...
- For a complete list of commands you can take a look to the UEFI shell specification:
 - http://www.uefi.org/sites/default/files/resources/UEFI_Shell_Spec_2_0.pdf

```

Shell> fs0:
fs0:\> dir
Directory of: fs0:\

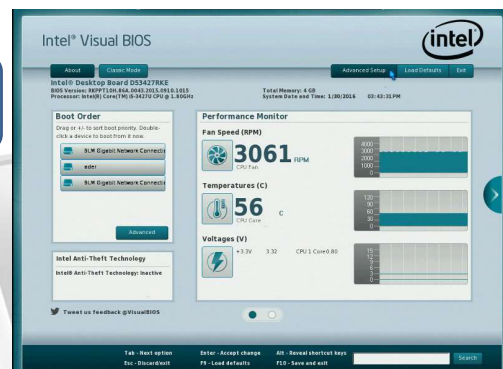
05/20/10  06:27p           5,628  SELViewer.txt
05/19/11  04:47p    <DIR>           8,192  \_images
05/20/10  06:20p      173,774  hrs.ini
05/20/10  06:27p      245,552  ipml.efi
05/20/10  06:27p         205  SEL.ini
05/20/10  06:27p      11,060  selenus.hlp
05/20/10  06:32p      98,356  selenus.str
05/20/10  06:32p     610,048  selview.efi
              7 File(s)    1,145,903 bytes
              1 Dir(s)
fs0:\> _
  
```

Stage 1: Hardware + UEFI (Boot Manager)



• UEFI Graphical User Interface

- Real world: graphical interface (the shell might still be present)



UEFI Boot Manager Configuration

- How can we modify Boot Manager configuration?

- For example, to modify boot order.

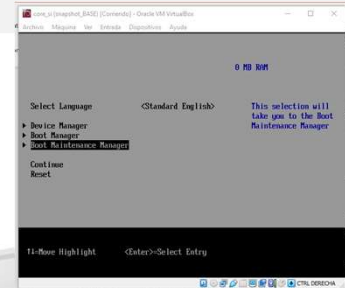
- Through the **UEFI graphical interface**

- Through the **UEFI shell app (bcfg command)**

- employed to modify NVRAM variables.
- `bcfg boot dump -v` (list current boot options)
- `bcfg boot add/rm/mv ...`



- Once the boot process is finished (kernel booted), through the **efibootmgr** command

- `# efibootmgr -v` (list boot entries)
- `# efibootmgr -c -d /dev/sda -p 2 -L "prueba" -l "\EFI\prueba\grubx64.efi"`
 - `(-c)` to create a new entry.
 - `(-d)` disk on which the EFI system Partition is hosted
 - `(-p)` partition number on which the EFI system Partition is hosted.
 - `(-L)` Label to use as boot entry.
 - `(-l)` path of the EFI image to boot.




Index

- Introduction
- Booting, Stage 1: Hardware+UEFI
- Booting, Stage 2: Bootloader (GRUB)**
- Booting, Stage 3: Kernel
- Booting, Stage 4: Systemd
- Shutting Down



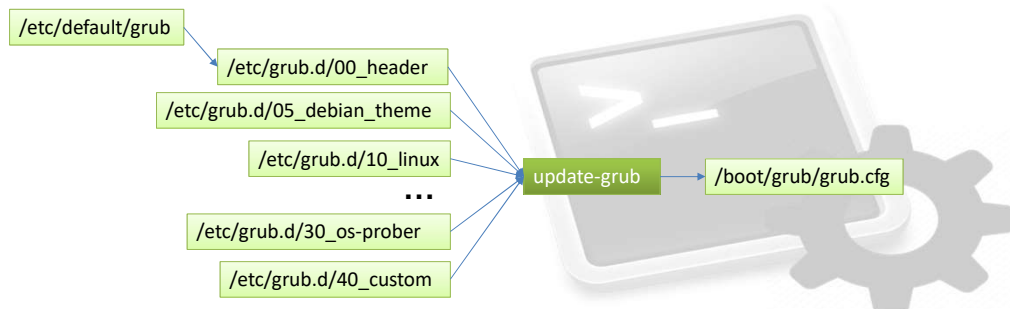
Stage 2: Bootloader (GRUB)

- "Piece" of software between UEFI and OS Kernel (during boot process)
 - Identify and load an appropriate OS kernel.
 - Provide configuration arguments for the kernel (-s for single-user mode).
- Task: Load Kernel into main memory to continue the boot process.
- **GRUB: GRand Unified Bootloader**
 - Developed by GNU project, default on most Linux distributions.
 - Two development branches: GRUB / GRUB2
 - Currently the most employed is GRUB2 (the one described in these slides).
 - Can work with file systems (ext2, ext3, ext4,...), directly accessing partitions.
 - Which is its current utility? multi-boot systems.
 - <https://www.gnu.org/software/grub/manual/grub/grub.html>



Stage 2: Bootloader (GRUB)

- **GRUB Configuration:**
 - Parameters such as: kernel to boot, boot options, boot password, etc.
 - GRUB reads its configuration params from **/boot/grub/grub.cfg**
 - grub.cfg should not be edited manually, created through **update-grub** command.
 - update-grub generates grub.cfg according to /etc/default/grub and /etc/grub.d/



Stage 2: Bootloader (GRUB)

- **GRUB Configuration:**
- The **/etc/default/grub** file:
 - GRUB_DEFAULT: default menu entry by menu position (starting at 0).
 - GRUB_CMDLINE_LINUX: params to add to the end of the “linux” command line.
 - GRUB_TIMEOUT: seconds to display menu before autoboot.
- The **/etc/grub.d/** directory (read during execution of update-grub):
 - 00_header: sets environmental variables (system file locations, video settings, etc.) and import preferences stored in /etc/default/grub.
 - 05_debian_theme: GRUB2 appearance (colours, background image, etc.)
 - 10_linux: identify kernels on the root device for the OS in use and creates entries (including associated recovery mode).
 - 30_os-prober: search for other OS
 - 40_custom: template for adding custom menu entries.



Stage 2: Bootloader (GRUB)

- **GRUB Command Line:**
 - GRUB supports a command-line interface for editing config file entries on the fly at boot time.
 - You can also boot non-listed OS, display system info and perform filesystem testing.
 - Some useful commands:
 - **boot**: boot the OS that was last loaded
 - **linux <path-to-kernel> [options]**: load a linux kernel
 - **reboot**: reboot the system
 - complete list: https://www.gnu.org/software/grub/manual/grub/html_node/Command_002dline-and-menu-entry-commands.html
 - Modifying kernel boot options at boot time:
 - **root=<path>**: root filesystem.
 - **ro/rw**: Mount root device read-only or read-write on boot.
 - **quiet**: Dissable most log messages.
 - **init=<command>**: Run specified binary instead of /sbin/init as init process.
 - **S**: boot kernel in single mode.
 - Complete list: <https://www.kernel.org/doc/html/v4.14/admin-guide/kernel-parameters.html>



Stage 2: Bootloader (GRUB)

- Having physical access to a system, stages 1 & 2 are a serious weakness.
 - Modifying boot options we could easily obtain superuser privileges.
- Increase protection a little bit: GRUB2 with password (still weak)
 - You can design a superuser with password to access GRUB shell.
 - Edit `/etc/grub.d/00_header` and at the end of the file add (remember to perform `update-grub` after that):

```
cat << EOF
set superusers="alumno"
password alumno 1234
export superusers
EOF
```


- Notice that password is in plain text. to encrypt it: `grub-mkpasswd-pbkdf2`

```
ubuntu@ubuntu:~$ grub-mkpasswd-pbkdf2
Enter password:
Reenter password:
Your PBKDF2 is grub.pbkdf2.sha512.10000.FC583738CA15A797C418C1EA7FFB85B9A21798D94B007BF5A57
9449728ADF249EABE1511C7B4277CB354092C0568E9008C304384D23F7B62F767.E657080F51EC8DE44B7053122
13BA9B59B1290013B92B68DAED9B45462E109F40CA6A935C263A4D87575302FF368036B4D73321DFC566C5697CA
```

```
cat << EOF
set superusers="alumno"
password_pbkdf2 alumno grub.pbkdf2.sha512.10000.FC58...
export superusers
EOF
```


Index

- Introduction
- Booting, Stage 1: Hardware+UEFI
- Booting, Stage 2: Bootloader (GRUB)
- **Booting, Stage 3: Kernel**
- Booting, Stage 4: Systemd
- Shutting Down



Stage 3: Loading the Kernel

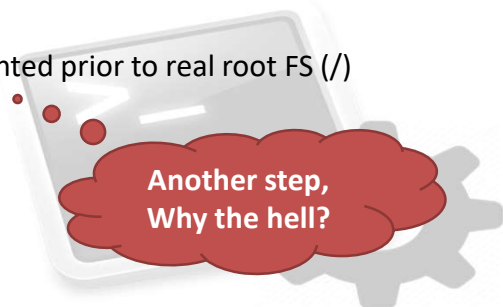
- The bootloader has loaded in memory kernel & ramdisk files
 - vmlinuz-4.9.0-4-amd64
 - initrd.img-4.9-0-4-amd64
- Once finalized Stage 2, kernel execution starts:
 - The Kernel **un-compresses** itself.
 - Detects memory map, the **CPU** and its features supported.
 - Starts the **display** (console) to show information through the screen.
 - Checks the **PCI bus**, creating a table with the peripheral detected.
 - Initializes the system in charge of **virtual memory** management, including swapper.
 - Initializes the **drivers** for the peripherals detected (Monolithic or modular).
 - Mount **file system** root ("/").
 - Calls the **systemd** process (Stage 4): PID 1, father of the rest of processes.



Stage 3: The init Ramdisk

- **RAM Disk:** fraction of main memory (RAM) formatted with a file system (tmpfs/ramfs)
 - Fast storage, but volatile!!
 - It is easy to create your own ramdisk...

```
# mkdir /tmp/ramdisk
# mount -t tmpfs -o size=1G myramdisk /tmp/ramdisk
```
 - That's it (everything will be lost when unmounting)
 - You can test write/read speed with dd command.
- **Initial RAM Disk (initrd):** transient root FS mounted prior to real root FS (/)



Stage 3: The init Ramdisk


- **RAM Disk:** fraction of main memory (RAM) formatted with a file system (tmpfs/ramfs)
 - Fast storage, but volatile!!
 - It is easy to create your own ramdisk...

```
# mkdir /tmp/ramdisk
# mount -t tmpfs -o size=1G myramdisk /tmp/ramdisk
```
 - That's it (everything will be lost when unmounting)
 - You can test write/read speed with dd command.
- **Initial RAM Disk (initrd):** transient root FS mounted prior to real root FS (/)
 - Main target: load the modules (HD drive) required to make the real FS available.
 - /sbin/init is executed
 - it mounts the “real” root file system
 - exec the /sbin/init on the “real” file system
 - after that, initrd file system is removed. The Kernel **un-compresses** itself.
 - Format: compressed cpio image

```
# cp /boot/initrd.img-XXX /tmp/initrd.img.gz
# gunzip initrd.img.gz
# cpio -i -make-directories < initrd.img
```

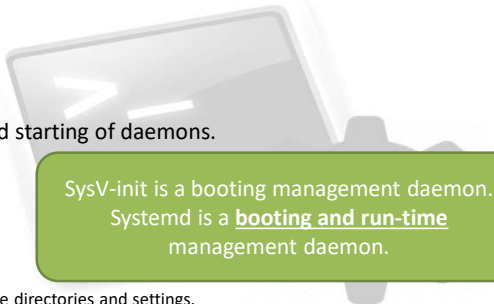
Index

- Introduction
- Booting, Stage 1: Hardware+UEFI
- Booting, Stage 2: Bootloader (GRUB)
- Booting, Stage 3: Kernel
- **Booting, Stage 4: Systemd**
- Shutting Down



Stage 4: Systemd

- Once the kernel is loaded, run the **system management daemon (/sbin/init)**.
 - Main goal: make sure the system runs the right **group of services (mode)** and daemons at any given time.
 - It provides a system and service manager running as PID 1.
 - In charge of startup tasks such as: setting computer name and time zone, check disk status, mount filesystems, configure network interfaces,...
 - Alternative implementations: SysV-init, BSD-init, **systemd**.
- Which are systemd main features?
 - aggressive parallelization capabilities (faster startup).
 - Socket and D-Bus activation for starting services, on-demand starting of daemons.
 - keeps track of processes using Linux control groups.
 - maintains mount and automount points.
 - logging daemon.
 - utilities to control basic system configuration.
 - hostname, date, locale, list of logged-in users, system accounts, runtime directories and settings,



SysV-init is a booting management daemon.
Systemd is a **booting and run-time** management daemon.

Stage 4: Systemd

- Systemd **Unit**: encoded information/configuration for any resource/service managed by systemd.
 - This is the primary object that the systemd tools know how to deal with.
- Available Systemd unit types:
 - **.service**: A system service.
 - **.target**: A group of systemd units (for startup). Involves in boot process.
 - **.automount**: A file system automount point.
 - **.device**: A device file recognized by the kernel.
 - **.mount**: A file system mount point. (alternative to fstab entry, as seen in
 - **.socket**: An inter-process communication socket.
 - **.swap**: A swap device or a swap file.
 - **.timer**: A systemd timer.
 - ...
- The behavior of each unit is defined and configured by its **unit file**.

Stage 4: Systemd

- **Unit file syntax:** (rsync daemon).
 - Internal structure organized with sections, denoted as: [section_name].
 - At each section, behavior is defined through key-value directives (one per line).

```
[Unit]
Description=fast remote file copy program daemon
ConditionPathExists=/etc/rsyncd.conf

[Service]
ExecStart=/usr/bin/rsync --daemon --no-detach

[Install]
WantedBy=multi-user.target
```

- Location of the Unit files:
 - /lib/systemd/system/, /etc/systemd/system/

Stage 4: Systemd

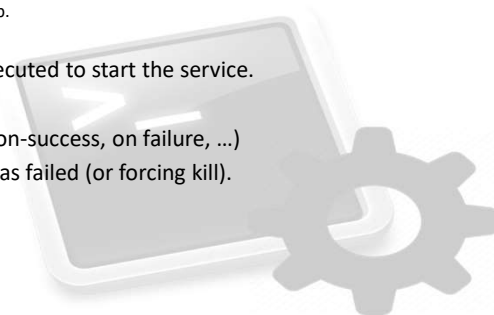
***Read man pages for a complete directive list**

- **[Unit] section directives***: define unit metadata and the relation to other units.
 - **Description/Documentation**: describe name, basic functionality and location for documentation about the unit.
 - **Requires**: lists units upon which current unit depends. If the current unit is activated, the units listed here must be activated (else this unit fails).
 - Requires directive can be replaced by creating a unit-file.requires dir in /etc/systemd/system and adding symlinks there to other unit files.
 - **Wants**: Similar to Requires but less strict (no activation required for the units listed)
 - Wants directive can be replaced by creating a unit-file.wants dir in /etc/systemd/system and adding symlinks there to other unit files.
 - **Requisite, Binds To, PartOf, Conflicts***
 - **Before**: The units listed will not be started until the current unit is marked as started.
 - **After**: The units listed will be started before starting the current one.
- **[Install] section directives***: define behavior of a unit if it is enabled or disabled (systemctl enable).
 - **WantedBy**: specify a dependency in a similar way to “Wants”. When a unit with this directive is enabled, the directory /etc/systemd/system/[unit].wants is created, with a symbolic link inside to create the dependency.
 - **RequiredBy**

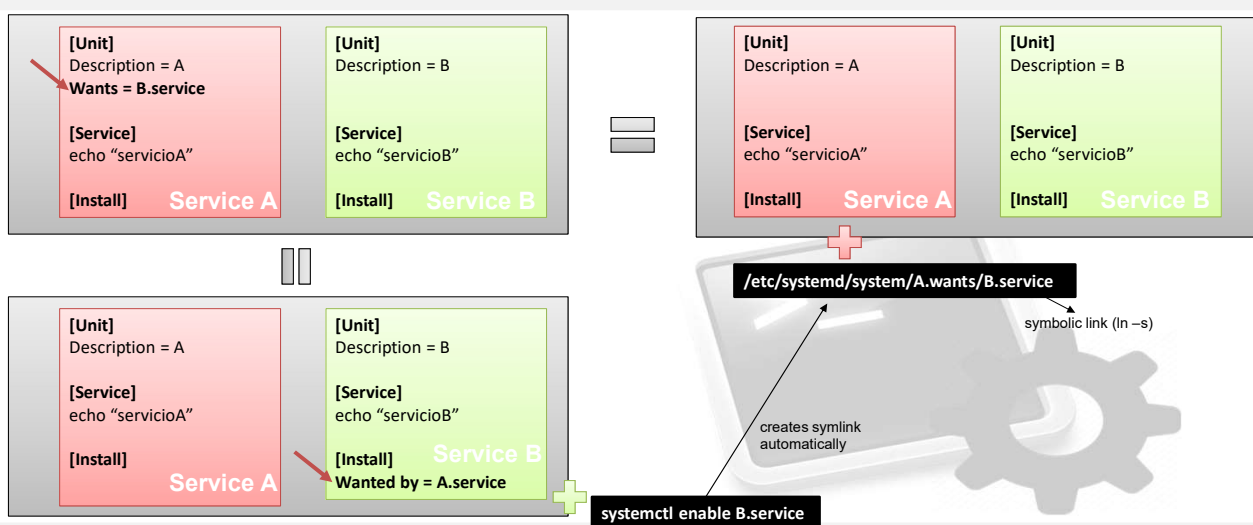
Stage 4: Systemd

*Read man pages for a complete directive list

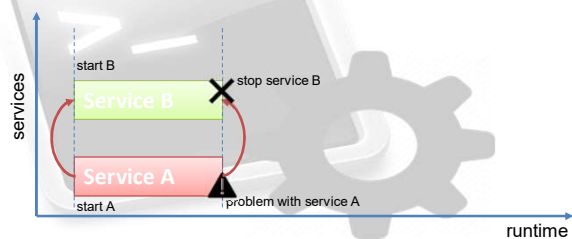
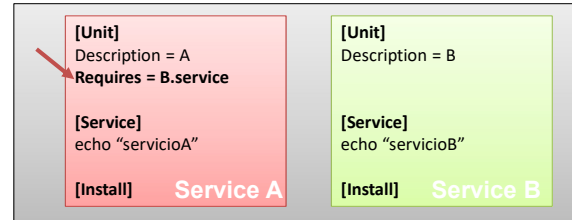
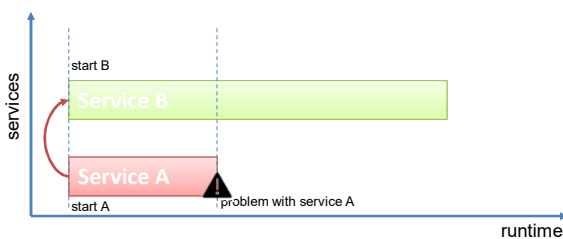
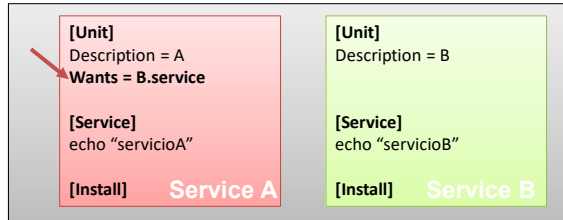
- **[Service]** section directives*: configuration of services.
 - **Type**: categorize service by their process and daemonizing behavior
 - simple: default type.
 - forking: the service forks a child process.
 - oneshot: systemd should wait for the process to exit before continuing on with other units.
 - dbus: the unit will take a name on the D-Bus bus.
 - notify: service will issue a notification when it has finished starting up.
 - idle: service will not run until all jobs are dispatched.
 - **ExecStart**: path and arguments of the command to be executed to start the service.
 - **ExecStop**: command needed to stop the service.
 - **Restart**: conditions to attempt automatic restart (always, on-success, on failure, ...)
 - **TimeoutSec**: when stopping, time to wait before marking as failed (or forcing kill).
 - ...




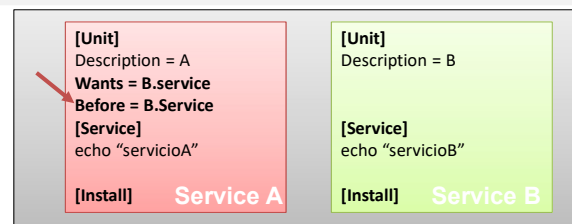

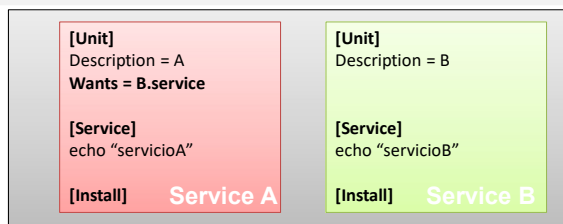
Unit File Relations (Definition)



Unit File Relations (Strength)



Unit File Relations (Timing)



Stage 4: Systemd

- Systemd **boot process** :

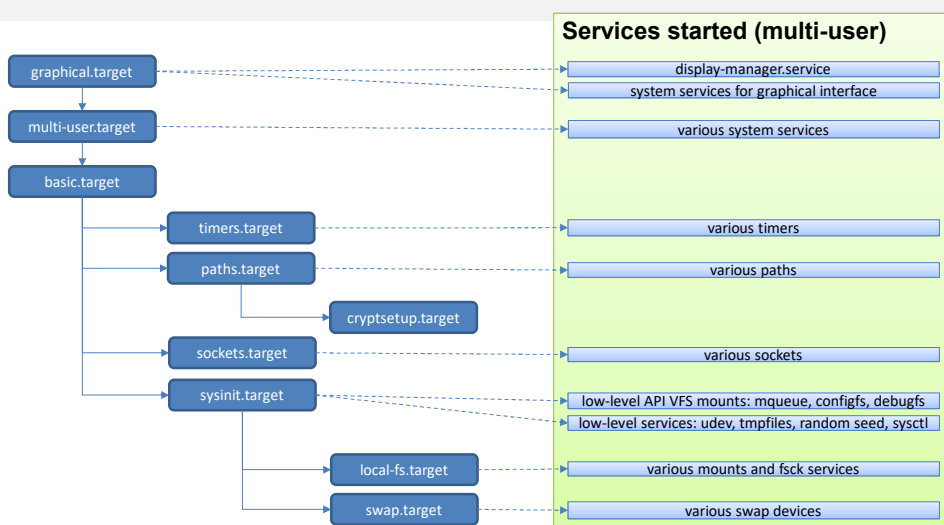
- boot & service management handled through **Targets**: special units employed to group boot units and start up synchronization processes.
- Targets are equivalent to SysV runlevels (see appendix for details), defining different operation modes.

runlevel	target	Description
0	poweroff.target	System halt
emergency	emergency.target	Bare-bones shell for system recovery
1,s, single	rescue.target	single-user mode
2-4	multi-user.target	Multi-user modes
5	graphical.target	Multi-user + GUI
6	reboot.target	system reboot

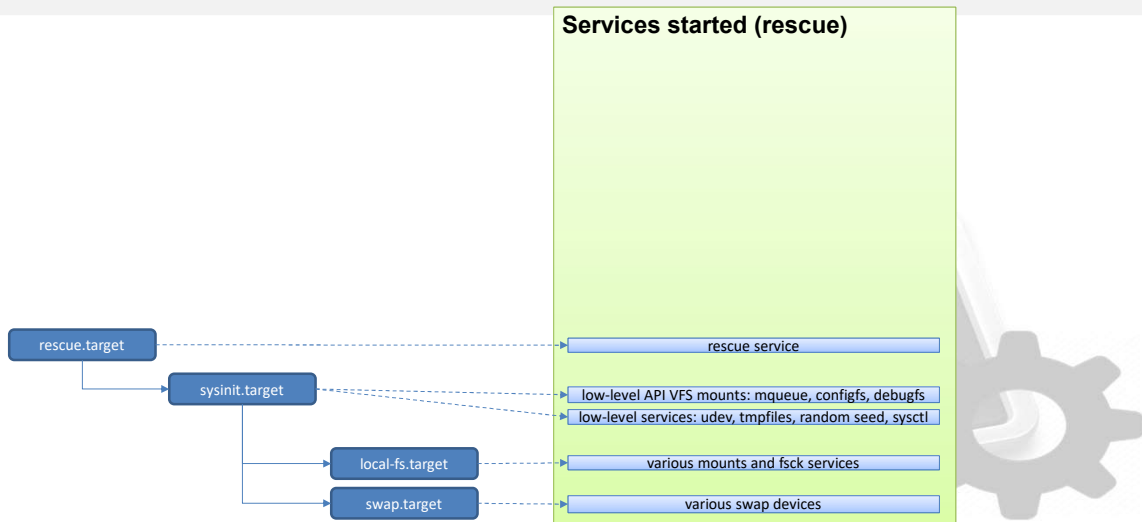
- Target relations to other units (targets or services) define the group of services started for each operation mode

```
[Unit]
Description=foo boot target
Requires=multi-user.target
Wants=foobar.service
After=rescue.service
```

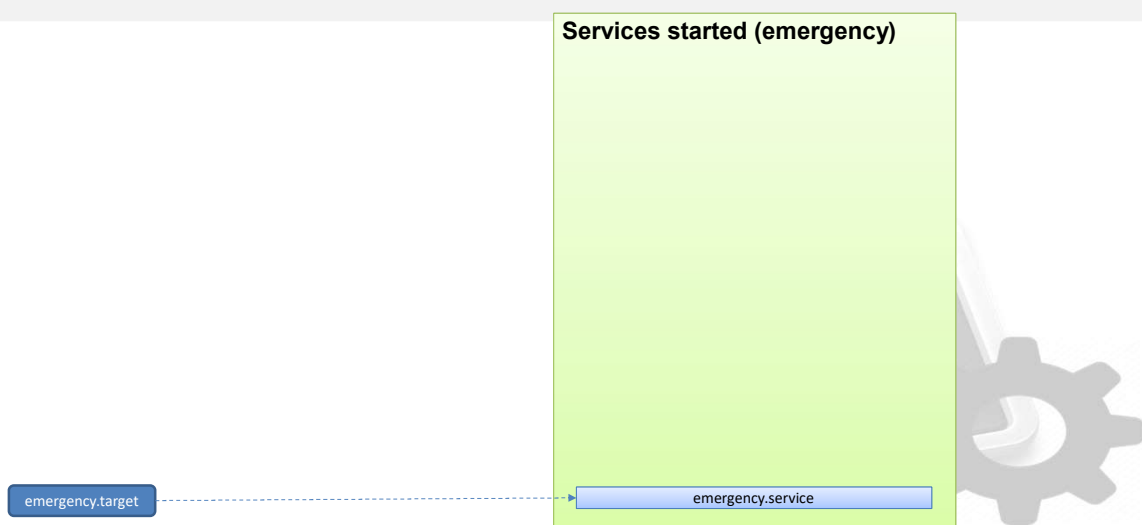
Stage 4: Systemd



Stage 4: Systemd




Stage 4: Systemd




Stage 4: Systemd

- Managing systemd through **systemctl** command
- Service administration:
 - `# systemctl list-unit-files [--type=service]`: list units available (of each type).
 - `# systemctl list-units`: list units active/waiting/failed.
 - `# systemctl status -l <unit>`: display service detailed information.
 - `# systemctl enable/disable`: activate/deactivate unit at boot.
 - `# systemctl start/stop/restart`: activate/deactivate/restart unit immediately.
- Operation mode administration:
 - `# systemctl isolate target`: change operation mode to target.
 - `# systemctl get-default`: see the target the system boots into.
 - `# systemctl set-default target`: change default target.
- System & Boot performance statistics:
 - `# systemd-analyze`.




Stage 4: Systemd

- Creating your own local services (do it in `/etc/systemd/system`)
 1. Create your own unit file (adapt one from `/lib/systemd/system`). Use the man pages to see a complete list of `[service]` and `[unit]` directives.
 2. Manage your dependencies,
 - Explicit: `Wants`, `Requires` directives inside unit file
 - Create the directories `unit-file.wants` or `unit-file.requires` in `/etc/systemd/system` and add symlinks there to other unit files.
 - Use the command `systemctl`: `# systemctl add-wants multi-user.target custom.service`
 3. Activate the dependencies: `# systemctl enable custom.service`



Index

- Introduction
- Booting, Stage 1: Hardware
- Booting, Stage 2: Bootloader
 - LILO
 - GRUB
- Booting, Stage 1+2 (UEFI)
- Booting, Stage 3: Kernel
- Booting, Stage 4: INIT
- **Shutting Down**




Shutting Down

- Never shut down directly (reset!).
 - If this rule is not respected, there is a high probability of losing or corrupting system files (with a bit of bad luck, fully broken system)
 - Intermediate Buffers for disk read/write. Synchronization.
- Never shut down without warning all system users
 - Periodically programmed shut-downs.
- Steps for a correct shut down:
 - Warn all the users previously.
 - Stop all the services associated to the target
 - Send the specific signal to all the processes to end their execution.
 - Users and processes still present, killed.
 - Subsystems shut down sequentially.
 - File System unmounted (synchronizes pending changes with disk)



Shutting Down

- Command **shutdown**:
 - Format: `/sbin/shutdown -<options> time message`
 - Option -r: reboot instead power off
 - Option -h: stop the system(with ACPI).
 - Message: message sent to all users.
 - time: delay to begin the shutdown (mandatory)
 - Format: hh:mm
 - Supports now+,minutes
- `/etc/shutdown.allow` or `inittab`
 - Avoid Ctrl+Alt+Del
- Other commands: `/sbin/halt`, `/sbin/poweroff`



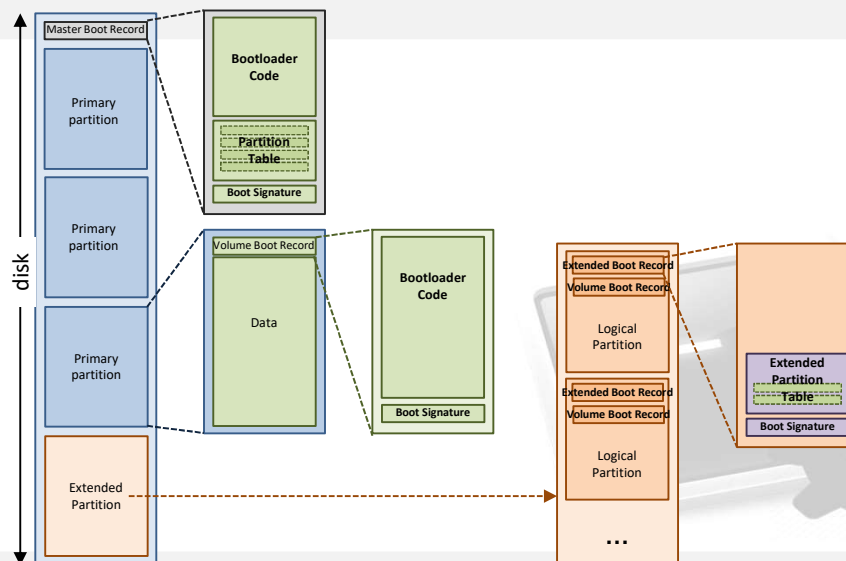
APPENDIX



BIOS Firmware

- **BIOS** (Basic Input/Output System):
 - 1975: First appearance in the Operating System CP/M.
 - It runs in real address mode (16 bit): 1MB of addressable memory.
 - 1990: appears “BIOS setup utility”: allows the user to define some configuration options (boot priority).
 - ROM customized for a particular HW. Provides a small library with I/O functions to work with peripherals (keyboard, screen). Very slow (protected to real mode).
 - Emerging applications require more and more BIOS support: Security, Temperature/Power metrics (ACPI), Virtualization extensions, Turbo-Boost ... (Hard to put all that in 1MB).
 - 2002: Intel develops an alternative firmware: EFI (/UEFI).

MBR Disks & Partitions



MBR Disks & Partitions

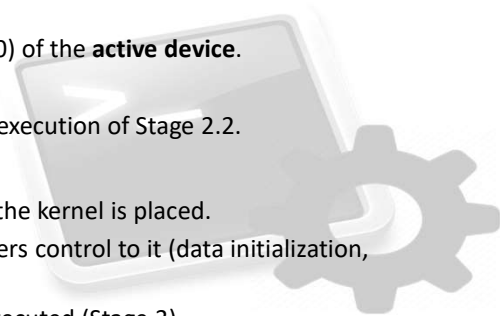
- **Master Boot Record (MBR):**
 - First block of the Disk, 512 Bytes.
 - **Partition Table:** information about four primary partitions: begin and end blocks, size, etc. (64 bytes)
 - **Boot Signature:** Numerical value indicating the presence of valid bootloader code in the code field (0x55AA) (2 bytes).
- **Volume Boot Record (VBR):**
 - First block of each primary partition.
 - Could contain bootloader code (indicated by Boot Signature).
- **Extended Partition:**
 - Partition that can be sub-divided into multiple **logical partitions**.
 - **Extended Boot Record (EBR):** First block of each logical partition. It only contains a partition table with two fields. **Extended partition table** forms a linked list with all logical partitions.

MBR Disks & Partitions

- **Linux Naming Convention:**
 - Remember: I/O devices are treated as files. Under directory /dev we find all system disks.
 - generic PC: 2 IDE controllers, each can have two devices (master/slave).
 - /dev/**hda**: first device (master) of the first IDE controller.
 - /dev/**hdb**: second device (slave) of the first IDE controller.
 - /dev/**hdc**: first device of the second controller.
 - /dev/**hdd**: second device of the second controller.
 - In a disk, each **primary partition** is identified with a number from 1 to 4.
 - /dev/**hda1**: first primary partition of the hda disk.
 - **Lógica partitions** start from 5.
 - /dev/**hda5**: first logical partition of hda disk.
 - In **SCSI devices** same naming convention, changing “sd” by “hd”
 - /dev/**sda1**

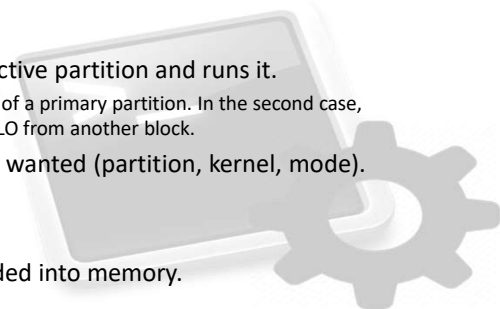
Bootloader in MBR

- Hardware requires an OS in charge of providing all the functionality in a computer.
- Target: loading in memory OS kernel and start running it. Loader with different locations: USB, CD, Disk ...
- **Stage 2.1:**
 - Located in **MBR**: 512 first bytes (block 0) of the **active device**.
 - Loaded in memory by BIOS (Stage 1).
 - Triggers, when executed, the load and execution of Stage 2.2.
- **Stage 2.2:**
 - Located in the **active partition**, where the kernel is placed.
 - Loads the kernel in memory and transfers control to it (data initialization, drivers, check CPU, etc.)
 - After this process, the **init** process is executed (Stage 3)



LILLO

- **Linux Loader:**
 - Two stage Bootloader.
 - Does not “understand” about operating system, neither about file system. Only works with physical locations.
 - Obsolete (but easy to follow for academic purpose)
- **Steps:**
 - Master boot loads LILLO from the first active partition and runs it.
 - LILLO can be in the MBR or in the Boot Block of a primary partition. In the second case, MBR contains the necessary code to load LILLO from another block.
 - LILLO requests the user the kind of boot wanted (partition, kernel, mode). Through a prompt.
 - LILLO loads the kernel and a ramdisk.
 - The kernel starts running once it is loaded into memory.



LILLO

- Configuration: `/etc/lilo.conf`

```
boot=/dev/hda #o by ID
map=/boot/map
install=/boot/boot.b
prompt
timeout=50
message=/boot/message
linear
default=linux

image=/boot/vmlinuz-2.6.2-2
    label=linux
    read-only
    root=/dev/hda2 #o by UUID
    initrd=/boot/initrd-2.4.2-2.img

other=/dev/hda1
    label=dos
    optional
```

Device where LILLO is installed
(IDE/SATA/Floppy...)

File with information about disk blocks
with the files required to boot system.

Loader Assembly code.

Kernel for booting and its options

Linux system partition (/). Not necessarily
a disk (usb loader).

Filesystem loaded in memory as a
ramdisk. Software support not provided
by the kernel to initialize the system.

Link to other loader (boot a different OS)

LILLO

- Configuration: `/etc/lilo.conf`

- Any change in the files employed in boot process (boot.b, kernel, ramdisk)
requires loader update:

- map file must reflect those changes, otherwise booting process is corrupted.
- Check if map file is updated: `# lilo -q`
- Update map file: `# lilo [-v]`

- A booting error cannot be fixed from the shell...

- Possible error sources:

- Installation of a new OS overwriting MBR (M\$)
- Failed kernel compilation
- Modification in boot files without map updating.

- Rescue Systems:

- mkbootdisk
- Installation Live CD (option rescue) or specialized (SystemRescueCD)

INIT (SysV)

- The init process performs the following tasks:
 - Step 1: **Configuration**: read from the file **/etc/inittab** the initial configuration of the system: Operation mode, runlevels, consoles,...
 - Step 2: **Initialization**: Runs the command **/etc/init.d/rc.S** (debian), which performs a basic initialization of the system.
 - Step 3: **Services**: According to the runlevel configured, runs the scripts/services pre-established for that runlevel.
- Runlevels (Operation modes)
 - Standard: 7 levels. Each distribution its own configuration (here Debian)
 - Level **S**: only executed at boot time (replaces **/etc/rc.boot**)
 - Level **0**: **Halt**. Employed to Shut down the system.
 - Level **1**: **Single User**. Maintenance tasks (no active network)
 - Level **2-5**: **Multuser**. All the network and Graphical services activated.
 - Level **6**: **Reboot**: Similar to level 0.

INIT (SysV)

- Step 1, Configuration. The file **/etc/inittab**:

```
# /etc/inittab: init(8) configuration.
# The default runlevel.
id:2:initdefault:

# Boot-time system configuration/initialization
# script. This is run first except when booting in
# emergency (-b) mode.
si::sysinit:/etc/init.d/rcS

# What to do in single-user mode.
~:S:wait:/sbin/sulogin

# /etc/init.d executes S and K scripts upon change
# of runlevel.
10:0:wait:/etc/init.d/rc 0
11:1:wait:/etc/init.d/rc 1
12:2:wait:/etc/init.d/rc 2
13:3:wait:/etc/init.d/rc 3
14:4:wait:/etc/init.d/rc 4
15:5:wait:/etc/init.d/rc 5
16:6:wait:/etc/init.d/rc 6
```

```
# Normally not reached, but fallback in case of
# emergency.
z6:6:respawn:/sbin/sulogin

# What to do when CTRL-ALT-DEL is pressed.
ca:12345:ctrlaltdel:/sbin/shutdown -t1 -a -r now
...

# Note that on most Debian systems tty7 is used by
# the X Window System, so if you want to add more
# getty's go ahead but skip tty7 if you run X.
1:2345:respawn:/sbin/getty 38400 tty1
2:23:respawn:/sbin/getty 38400 tty2
3:23:respawn:/sbin/getty 38400 tty3
4:23:respawn:/sbin/getty 38400 tty4
5:23:respawn:/sbin/getty 38400 tty5
6:23:respawn:/sbin/getty 38400 tty6
```

INIT (SysV)

- Step 1, Configuration. The file /etc/inittab:
 - Line format: **id:runlevels:action:process**
 - **id**: identifier for the entry inside inittab
 - **runlevels**: execution levels for that entry (empty means all)
 - **action**: What must init do with the process.
 - wait: wait until it finishes
 - off: ignore the entry (deactivated)
 - once: run only once
 - respawn: rerun the process if it dies
 - sysinit: ask the user what to do with that entry
 - Special: ctrlaltdel
 - **process**: sh line indicating init which process to start when this entry is reached.



INIT (SysV)

- Step 2, Initialization. The file /etc/init.d/rc:
 - Input parameters: the runlevel. Example rc 2: multiuser
 - Tasks:
 - Establish PATHs
 - Load swap space: swapon
 - Check and mount local filesystems (/etc/fstab)
 - Activate and configure the network
 - Remove not necessary files (/tmp)
 - Configure the kernel. Load modules: Drivers (managing dependencies)
 - Triggers the startup of the services associated to the runlevel.
 - Modifying the runlevel: command init, telinit
 - Allows changing from one runlevel to another
 - ¿Single User?
 - Restore original state.



INIT (SysV)

- Step 3, services. The directories `/etc/init.d` and `/etc/rcN.d`:
 - All the services available are found in `/etc/init.d`
 - Examples: `cron`, `ssh`, `lpd`, ...
 - How to indicate each runlevel which ones to start?
 - With a special directory, `/etc/rcN.d/` (being N the runlevel).
 - In these directories a list of links to the services is found.
 - The directory `/etc/rcN.d/`
 - The links begin with letters “S” or “K” plus two digits (execution order).
 - “S”: executed in ascending order when a runlevel is started (`ssh start`).
 - “K”: executed in descending order when shutting down (`ssh stop`).
 - These links are controlled with “`update-rc.d`”
 - `S99local`: script to perform local configurations
 - minor booting aspects: auxiliary kernel modules, personalized services,...
 - Employed by the administrator
 - It really runs the script `/etc/rc.local`



INIT (SysV)

- Step 3, services. The directories `/etc/init.d` and `/etc/rcN.d`:
 - The directory `/etc/rcN.d/`

```
pablo@si:/etc/rc2.d$ ls
README          S03cgroupfs-mount  S03vboxdrv       S05cups
S01bootlogs     S03cron            S04avahi-daemon  S05cups-browsed
S01syslog       S03dbus            S04docker        S05saned
S02apache2      S03exim4           S04lightdm
S06plymouth
...
```

```
pablo@si:/etc/rc6.d$ ls
K01alsa-utils   K01network-manager K02avahi-daemon  K06rpcbind
K01apache2      K01plymouth         K02vboxdrv       K07hwclock.sh
...
```



INIT (SysV)

- **Manual administration of services:**
 - After booting process, services can be modified (stop running services or start new services).
 - Directly through its script (example ssh):
 - # /etc/init.d/ssh [stop/start/restart/status]
 - Or through the command service:
 - service --status-all: reads /etc/init.d/ verifying service state [+] [-] [?]
 - These changes are volatile (lost after reboot).
 - Permanent with update.rc-d
 - Checking possible errors concerning boot process
 - # tail -f /var/log/messages (Another important files: syslog, daemon.log)
 - # ls -lart /var/log



INIT (SysV)

- **Manual administration of services:**
 - Examples of start script and services command:

```
#!/bin/sh
#SIMPLIFICADO
[ -f /usr/local/sbin/sshd2 ] || exit 0

PORT=

PORT=$(grep Port /etc/ssh2/sshd2_config | awk '{ x = $2 } END {print x}' -)
if [ "X$PORT" = "X" ]
then
    PORT=22
fi

# See how we were called.
case "$1" in
    start) # Start daemons.
        echo -n "Starting sshd2 in port $PORT: "
        /usr/local/sbin/sshd2
        echo "done."
        ;;
    stop) # Stop daemons.
        echo -n "Shutting down sshd2 in port $PORT: "
        kill `cat /var/run/sshd2_$PORT.pid`
        echo "done."
        ;;
    restart)
        $0 stop
        $0 start
        ;;
    *)
        echo "Usage: sshd2 {start|stop|restart}"
        exit 1
esac
exit 0
```

